


Endoscopic video clips (unified cross-orifice)

Hierarchical JEPA predictors (domain-conditioned)



Predict plannable representations, not pixels. Splits are video-level; sampling is domain-balanced.