

Cross-dataset h=4 latent forecast versus persistence retrieval (2/2, video-level validation)


Last observed frame

Ground-truth terminal future



Forecast-latent retrieval

Persistence-latent retrieval




STIR
h=4 cos 0.972 vs 0.874
h=4 win 100.0%






Hamlyn stereo laparoscopy
h=4 cos 0.987 vs 0.823
h=4 win 100.0%



Endoscopes
h=4 cos 0.975 vs 0.853
h=4 win 94.6%



EndoVis 2017
h=4 cos 0.974 vs 0.864
h=4 win 100.0%



EndoVis 2018
h=4 cos 0.974 vs 0.859
h=4 win 100.0%

